

Alpha-Factory v1: Multi-Agent AGENTIC α -AGI World Model Demo

MONTREAL.AI – AGI-Alpha-Agent-v0 Extension

April 25, 2025

1 Introduction and Objectives

The **Alpha-Factory v1** (α -Factory v1) demo showcases a **large-scale foundation world model** driven by a **constellation of autonomous agents**. The goal is to generate **diverse synthetic environments** and train **general, robust agents** on an open-ended curriculum, inching toward α -ASI (*artificial superintelligence*). This project builds on MONTREAL.AI’s existing **AGI-Alpha-Agent-v0** codebase and incorporates cutting-edge ideas in AI research. We aim for a **production-ready, flawless implementation** that can be deployed by non-technical users, demonstrating emergent general intelligence through multi-agent collaboration.

1.1 Key Objectives

- **Multi-Agent Orchestration:** Leverage at least five integrated agents from the Alpha-Factory suite (planner, learner, evaluator, environment-generator, etc.) working in concert.
- **Open-Ended World Generation:** Autonomously create and evolve diverse training environments (virtual worlds, tasks, simulations) to continually challenge and improve agents.
- **Advanced Training Loops:** Implement both *MuZero-style model-based learning* and *POET-style co-evolution* of environments and agents for robust skill acquisition.
- **Integration of AI Protocols:** Use **OpenAI’s Agents SDK**, **Google’s ADK**, **Agent2Agent (A2A) protocol**, and **Anthropic’s Model Context Protocol (MCP)** to enhance interoperability, security, and learning capabilities.
- **User-Friendly Deployment:** Provide a simple UI/REST API for non-experts, a CLI for developers, and containerized deployment (Docker/K8s) for easy scaling.
- **Antifragility & Robustness:** Design the system to become *more resilient under stress* (antifragile) and *secure by default*, with broad applicability across industries: *Outlearn*, *Outthink*, *Outdesign*, *Outstrategize*, *Outexecute* in any domain.

2 Architecture Overview

Alpha-Factory v1 is an **antifragile multi-agent architecture**. It consists of an orchestrator and a network of specialized agents, all built on the existing codebase’s patterns and extended for this demo. Each agent has a distinct role, and together they form an *agentic α -AGI* system where the whole is greater than the sum of its parts:

2.1 Orchestrator (Macro-Sentinel)

The central brain coordinating all agents. It spawns agents, assigns tasks, and manages the iterative training cycles. The orchestrator uses the **A2A protocol** for agent communication, enabling independent modules to share goals and state *regardless of framework or language*. It ensures that environment generation, learning, and evaluation proceed in sync, adjusting difficulty and focus as needed.

2.2 Environment Generator Agent

This agent creates **synthetic world models and tasks**. Drawing from **POET** (Paired Open-Ended Trailblazer) principles, it generates a **diverse and ever-expanding curriculum** of environments. Each environment can be a game level, a puzzle, a simulated physics world, or any scenario that challenges the agents. The generator uses *quality-diversity (QD)* and *open-endedness* algorithms to introduce novelty and complexity continually, ensuring the agent never outgrows its training data. As Jeff Clune’s *AI-GA* paradigm suggests, tasks and data are *learned and evolved* rather than hand-designed.

2.3 Learning/Planning Agent

A reinforcement learning agent that interacts with the environments to acquire skills. This agent utilizes a **MuZero-style approach** — learning an internal model (for reward, value, policy) and planning with MCTS (Monte Carlo Tree Search) *without needing a perfect simulator*. It can handle partial observability and stochastic worlds. Over time, it builds a **world model** to support lookahead planning and generalization across tasks.

2.4 Curriculum & Evaluation Agent

This agent monitors the performance of the learning agent on various environments and adjusts the curriculum accordingly. Inspired by **POET co-evolution** and Silver & Sutton’s “*Era of Experience*”, it ensures the training data evolves as the agent becomes stronger.¹ It might transfer the agent to harder versions of tasks, generate new challenges when old ones are mastered, or revisit simpler tasks if regressions are detected. **Curriculum learning** is automated: the agent “*learns from the environment’s feedback to continuously improve itself*”, rather than relying on human-designed lesson plans.

2.5 Knowledge/Memory Agent

Using **Model Context Protocol (MCP)**, this agent maintains and provides external context or memory to other agents. It can store learned skills, important world facts, or use an *LLM-powered tool* to summarize past experiences. As the learning agent explores, the knowledge agent may supply relevant background info or previously learned strategies via a secure standardized interface (like “USB-C” for plugging context into the world model).

2.6 Control & Safety Agent

Given the high level of autonomy, this agent watches for unsafe strategies or catastrophic failures. It enforces constraints (e.g., resource limits, ethical boundaries) and can intervene or reset an environment if needed. Drawing on best practices (from OpenAI’s *Practical Guide to Building Agents*), it ensures exploration remains safe, aligned, and robust against adversarial scenarios.

2.7 Interface Agent(s)

These agents handle communication with external users and systems. One sub-agent might offer a **REST API** and **UI dashboard** for monitoring training, while another exposes a **CLI** for power-users. They translate user intents into orchestrator actions and vice versa.

¹<https://www.techrepublic.com/article/news-ai-era-experience-silver-sutton/>

The design is modular: each agent can be replaced or upgraded without disrupting the others, thanks to standardized *A2A* and *MCP* protocols and the orchestrator’s clear interface.

3 Open-Ended Environment Generation

3.1 Diverse Synthetic Worlds

The environment generator agent employs techniques from *open-ended algorithms* like POET and *quality-diversity (QD) search*. It maintains a population of environment definitions and a population of agent solutions, co-evolving them:

- Start with a simple base environment and an untrained agent.
- Periodically mutate or perturb environments (new obstacles, parameters, or rules).
- If the agent can’t solve a new environment, shelve or adjust it; if solved, increase complexity further.
- Over time, an *ever-expanding tree* of tasks emerges, some of which are extremely hard (reserved for future).

To maintain **quality diversity**, novelty search and diversity metrics ensure the system avoids repetitive tasks and fosters broad skill coverage.

3.2 Curriculum and Experience Engine

Drawing from *Silver & Sutton’s “Era of Experience,”* the system treats experience as the teacher: the curriculum agent adapts the sequence of tasks based on agent performance:

- Fast-track the agent to tougher environments if it masters a task easily.
- Provide intermediate stepping-stone tasks or shaping rewards if the agent struggles.
- Revisit earlier tasks to avoid catastrophic forgetting and reinforce transfer across tasks.

This *algorithmic mentor* “learns how to teach” over time, potentially discovering training strategies beyond human intuition.

4 MuZero-Style Learning and Planning

Agents train within each environment using a **MuZero-style RL loop**:

- **Learned World Model:** A neural network learns value, reward, and policy predictions, plus a latent state representation useful for planning.
- **Monte Carlo Tree Search (MCTS):** For each decision, the agent performs lookahead in its learned model, simulating future action sequences.
- **No Hard-Coded Rules:** The agent infers environment dynamics by trial-and-error; it is never given explicit rules.
- **Training Loop:** Episodes are stored in a replay buffer and used to train the model. As new environments arise, they are added to the mix, pushing the agent to adapt.

This delivers strong in-environment performance and supports generalization via an internal predictive model that can handle many tasks.

5 POET-Style Co-Evolution Loop

Above the MuZero loop, a **POET-style** outer loop co-evolves tasks and agents:

1. **Environment Proposal:** The generator proposes new environments, ensuring they meet a minimal solvability criterion.
2. **Agent Adaptation:** The agent (or a copy) attempts the new environment; if it succeeds, that environment is “solved” and may spawn more complex variants.
3. **Incorporation and Transfer:** Skills learned transfer back into the agent’s main policy; new environments join the curriculum for ongoing practice.
4. **Iteration:** This loop continues indefinitely, expanding both agent capability and environment diversity.

By *co-evolving* problems and problem-solvers, the system fosters continuous, open-ended growth, inching toward α -ASI.

6 Integration of Advanced AI Frameworks

The system is **future-proof** via modular integration of several frameworks:

6.1 OpenAI Agents SDK

Enables multi-step reasoning agents using LLM-based planning. Optional if an API key is provided; otherwise, local models are used.

6.2 Google ADK (Agent Development Kit)

Standardizes agent packaging and lifecycle, facilitating microservice/container deployments and scaling across clusters.

6.3 Agent2Agent (A2A) Protocol

Defines a common message bus, letting agents share states, requests, and capabilities with well-defined *Agent Cards*.

6.4 Model Context Protocol (MCP)

Structures how models (LLMs or otherwise) receive context. Prevents prompt injection and ensures standardized data provision.

6.5 Best Practices: Safe & Effective Agents

We implement alignment, reward-hacking safeguards, and *Anthropic’s* constitutional AI principles as needed. Agents have well-defined roles/scopes, and any external calls or code executions are sandboxed.

7 User Interaction and Deployment

7.1 Web UI Dashboard

A lightweight web app (Streamlit, Flask, or similar) provides real-time monitoring and control of training and environment generation. Users can observe metrics, environment visuals, and agent performance.

7.2 REST API

A RESTful API (FastAPI or Flask) enables remote control: start/stop training, fetch logs, inject new tasks, etc. This allows integration with other systems or services.

7.3 Command-Line Interface (CLI)

Advanced users can run a headless training loop, do evaluations, or enter an interactive REPL session with the orchestrator. Logs and metrics are output to console/files.

7.4 Configuration & Extensibility

All agent behaviors and environment types can be configured via YAML or Python config objects. We include hooks for easy extension with new agent classes or environment domains.

7.5 Deployment Options

Local Python Execution Everything runs on a single machine with Python 3.x. No external services required.

Docker Container We provide a pre-built Docker image with all dependencies. Running a single container launches the orchestrator and UI.

Kubernetes (K8s) For large-scale or cloud deployments, we supply Helm charts/manifests to run each agent in its own pod, coordinated via the orchestrator.

Security and Isolation Agents are sandboxed, ensuring no malicious code execution or resource exhaustion. The REST API is token-protected, and logs/models can be encrypted at rest.

8 Ensuring Antifragility, Security, and Robust Intelligence

8.1 Antifragility

System stress (e.g., failing a difficult environment) triggers adaptive responses. Meta-learning modules or the curriculum agent can respond by breaking tasks into sub-steps or adjusting exploration.

8.2 Robustness

We test newly generated environments for validity and resilience. Agents are also subjected to noise and perturbations to ensure stable policies.

8.3 Security

We enforce strict role-based scopes: environment manipulation does not directly overwrite agent weights. External APIs are rate-limited, and content filters check any code output from an LLM agent.

8.4 Emergent α -ASI Behavior

Though true ASI is beyond one demo, we aim to show *glimmers of robust general intelligence*. Meta-learning, curriculum generation, and multi-agent synergy may yield surprising, creative solutions indicative of advanced problem-solving. All existing Alpha-Factory functionality remains and is extended, ensuring incremental progress toward α -ASI.

9 Conclusion and Deliverables

We deliver the **Alpha_ASI_World_Model** demo as a new directory in the repository, containing:

- **Codebase:** Production-level Python code with a clear agent architecture.
- **Documentation:** A Markdown user guide explaining setup, usage, and system design, plus agent orchestration/training loop diagrams.
- **Tests:** Automated tests validating environment generation, agent communication, and training.
- **Demo Scenarios:** Preset scenarios (Mini-Worlds, Physics Learning, Social Agents, etc.) illustrating open-endedness, planning, and multi-agent interaction.
- **Emergent Behavior Showcase:** Logs or a report highlighting interesting or creative solutions found during runs.

By integrating advanced frameworks, automated curriculum design, robust RL, and user-friendly deployment, **Alpha-Factory v1** establishes a **comprehensive foundation for α -AGI research**. It learns from open-ended experience, challenges itself continually, and bridges simulation with real-world applicability — stepping closer to true artificial superintelligence through **self-driven open-ended learning**.