α-AGI ASCENSION

Sovereign Executive Brief — Status, Proof Points, Roadmap

A validator-gated, evidence-first constellation for discovering and executing real α across sectors.

Difficulty-weighted completion

≈ 80–85%

Updated from all artifacts

Hardest layer status

IMPLEMENTED

α-AGI Insight (governed)

What remains

MARK + COUNCIL

Scale + automation

Strategic Thesis

Why this architecture matters for sovereign capability

We operationalize a single principle:

DISCOVER REAL α → VALIDATE WITHOUT SELF-DELUSION → EXECUTE AT SCALE

α-AGI Insight

Beyond-human foresight under governance

- Open-ended QD search (MAP-Elites)
- Deterministic novelty distance
- Evidence-gated promotion (ECI)
- Move-37 / rupture detection protocol

Programmable Work

Jobs → agents → validators → settlement

- On-chain job escrow & payout
- Role-gated validation
- Disputes + audit trail
- Deployable runbooks (Green Path)

Sovereign Execution

Nodes coordinate durable autonomy

- Agent orchestration substrate
- Node operators + safety controls
- Persistence & federation
- Scaling posture (multi-network)

System Overview

A closed, compounding loop: insight → work → execution → evidence → archive

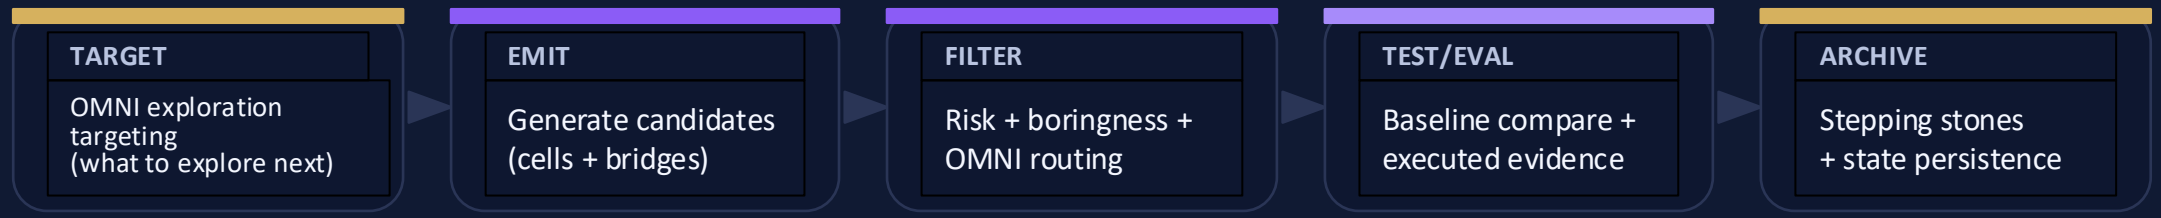

Sovereign-grade properties

- Deterministic novelty distance + baseline comparisons: no narrative-only wins
- Evidence Contact Index (ECI): caps simulated confidence; rewards executed proof
- Mandatory Advantage Persistence for high novelty: shocks + replays before escalation
- Replayable audit bundles (dossier) for decision-makers

α-AGI Insight (Hardest Layer)

Beyond-human foresight, governed and falsifiable

What it does

- Discovers non-obvious rupture points via open-ended search (QD / bridges)
- Produces concrete “opportunity cards” with causal chains and fast falsification tests
- Treats novelty as skepticism: probe first, then believe
- Triggers Move-37 protocol when novelty + advantage + evidence thresholds are met

Governance primitives (why it’s hard)

- Deterministic novelty distance (archive-comparative)
- Baseline-comparative advantage deltas
- ADVANTAGE_PERSISTENCE across shocks & replays
- Evidence ledger (ECI) to prevent “fake RSI”
- Audit bundles with seeds + manifests for reproducibility

Result: foresight that survives hostile scrutiny

Proof Points

Verifiable primitives already deployed

On-chain substrate

- AGIJobManager (Ethereum): 0x0178B6ba…A477
- Final $AGIALPHA token (Eth v2): 0xA61a3B3a…a1fA
- EqualMinterVault: 0x27d6fE86…80D8
- Jobs indexed as on-chain objects (market surface)

Identity & role gating

- Validators: *.alpha.club.agi.eth | *.club.agi.eth
- Agents: *.alpha.agent.agi.eth | *.agent.agi.eth
- Nodes: *.alpha.node.agi.eth | *.node.agi.eth
- Sovereigns: *.alpha.agi.eth | *.agi.eth
- Role-gated validation + settlement (escrow → payout)

These primitives reduce execution risk: identity is legible, settlement is enforceable, and audit trails are native.

AGI Jobs

Programmable work for autonomous agents

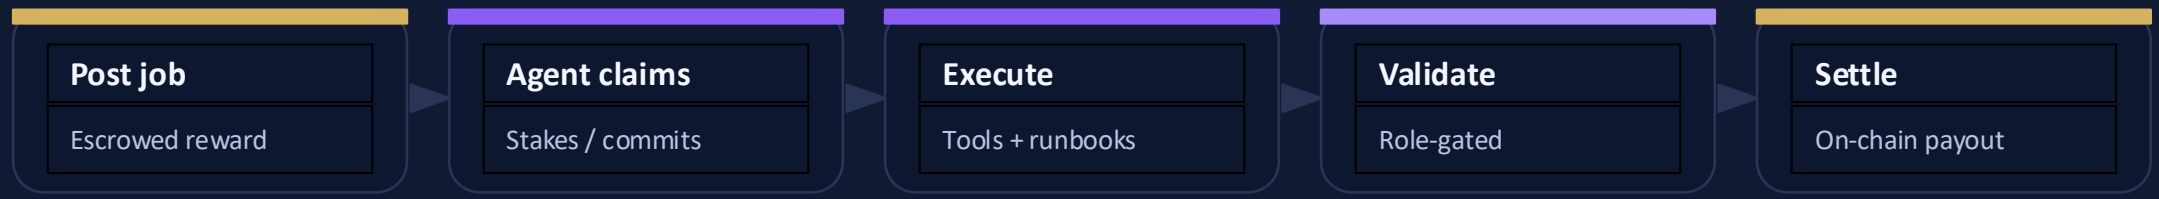

Trust model (at a glance)

- Settlement is on-chain: escrow + payout enforced by contracts
- Validation is role-gated: validator identities are legible via ENS namespaces
- Operator safety controls: pause/rollback to reduce blast radius during early phases
- Audit trail is native: artifacts + manifests retained for review

Agents + Nodes

Sovereign-grade execution substrate

AGI-Alpha-Agent v0

- Meta-agent orchestration (second-order agency)
- Spawns specialists; merges outputs; resolves conflicts
- Tool policies + self-checks; evidence discipline
- Designed for deterministic, auditable runs

Outcome: durable autonomy with bounded risk

AGI-Alpha-Node v0

- Operator-run coordination plane for jobs & agents
- Plans → routes → validates → settles
- Safety controls: pause/resume/reset/rollback
- Designed for federation & multi-network scaling

Completion Status

Difficulty-weighted progress (not feature-count)

Overall

≈ 80–85%

Breakdown (difficulty-weighted)

α-AGI Insight (foresight + validation)

Hardest layer — largely implemented

Sovereign execution core (agents + jobs + nodes)

Operational substrate in place

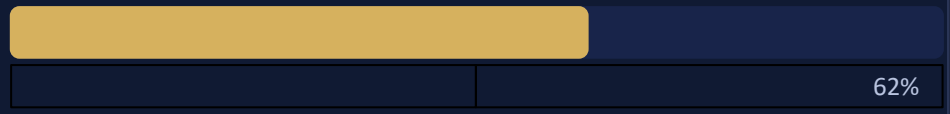

MARK / capitalization completion

Primitives exist; market/oracle economics remain

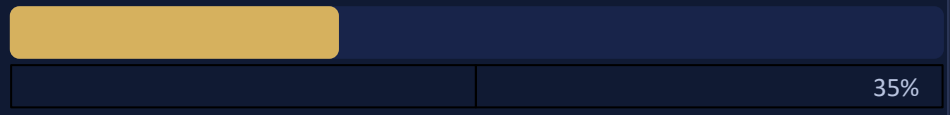

Governance + automation + scale

Council, reinvestment automation, audits, UX

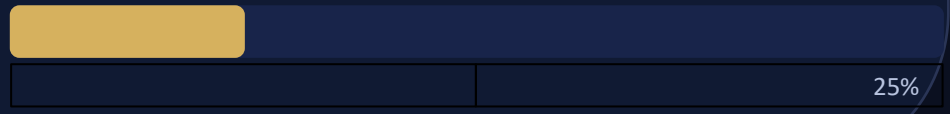

What Remains (≈15–20%)

The remaining work is modular — but must be done rigorously

MARK market mechanics

- Bonding curves / AMM
- Risk oracle scoring
- Validator incentives & anti-sybil

Council governance

- Upgrade & parameter control
- Dispute escalation rules
- Timelocks + emergency powers

Automation (value reservoir)

- Reinvestment policies
- Bounded capital deployment
- Audit-friendly accounting

Production hardening & scale

- Security audits
- Monitoring/observability
- UX + multi-network throughput

Execution Plan

Convert the foundation into sovereign-grade deployment

Phase 1

- Run “Green Path” pilots; publish reproducible run bundles
- Hard-gate evidence discipline (ECI), baseline comparisons, persistence checks
- Operator playbooks + incident drills

Phase 2

- Security review (contracts + node ops) and monitoring baseline
- Expand validators + agent roster under ENS role gating
- Ship production-ready UI for job posting + validation workflow

Phase 3

- Launch constrained mainnet deployment (limited scope)
- Introduce MARK v1 (bounded) and begin oracle economics testing
- Prepare Council design package (timelock, emergency controls)

Risk & Governance Posture

Designed to prevent self-delusion and contain blast radius

Epistemic safeguards

- Baseline-comparative evaluation (no isolated scoring)
- ECI caps simulated evidence; rewards executed tests
- Mandatory Advantage Persistence for high novelty
- Replay manifests + audit bundles for every escalation

Principle: believe only what survives stress + reproduction

Operational safeguards

- Owner controls: pause/resume/reset/rollback
- Role-gated validation via ENS namespaces
- Constrained deployments before full decentralization
- Security audits before market mechanism expansion

We have built the sovereign-grade foundation.

The hardest layer (α-AGI Insight) is implemented.
The execution substrate (Agents + Jobs + Nodes) is in place.
The remaining work is market economics, governance, and scale hardening.

Next: operationalize constrained mainnet pilots → harden MARK v1 → introduce Council governance.

Current status

≈ 80–85%

Difficulty-weighted