AGI ALPHA FLYWHEEL — When Capability Creates Opportunity

Transformational Intelligence, Experience-Trained Specialists and the Expanding Frontier of Profitable Work

Vincent Boucher — President of MONTREAL.AI
Executive brief · Additive opportunity-discovery edition v4 · 27 October 2026

FIND YOUR ALPHA. BUILD YOUR SUCCESSOR Ω.

The result

We constructed a reusable manufacturing inspection-evidence instrument and transferred it to a fresh recipient with new mission checks and current local research authority. It leaves the enterprise better equipped to trace claims, find counterexamples and specify missing evidence. **This edition does not establish a profitable new market, measured Operational Alpha, or commercial compounding.** Its strongest commercial result is a prospective decision to preserve cash and use the existing expert alternative.

The selected opportunity is consequential and bounded: a receiving-quality engineer needs to know whether another provider's inspection evidence supports its obligations, and a successor needs an exact evidence-acquisition plan for investigation. Existing digital-thread tools are strong alternatives. The new instrument addresses an inspectable cross-provider semantic problem, not the invention of STEP, QIF, geometry matching or inspection automation.

Exploration changed the direction

Ten candidate directions and challenges are preserved. The team investigated software failures, scientific dependency/evaluation problems, causal evidence, robotic rheology, manufacturing qualification, thermal metrology and laboratory portability before selecting the flagship. The familiar microgrid work remains a historical fixture.

A particularly useful negative pilot reproduced an old Dask/Zarr ZIP corruption report but found that current software passed. That proposed repair opportunity was rejected rather than promoted into a large simulator. Thermal examples also lost their proposed gap when known symmetry and manufacturing information were included. The selected manufacturing direction survived because actual NIST/MTC files supplied both a complete geometry/result trace and concrete mismatches that an audit must resolve.

What was built

The instrument reads exact QIF 2.1 sources, preserves their literal units, tolerances and statuses, constructs geometric witnesses from actual AP242 CAD, checks interval implications and recorded values, and binds repair requests to specific source obligations. It distinguishes numerical records, physical families, individual features and part release.

Plate development imported a valid solid with **261 faces** and constructed **17 requirement–feature–geometry–result traces**. A recorded diameter of **3.017 mm** passes its local interval while the containing inspection remains **FAIL**. That distinction survives the report and transfer.

Exact QIF 2.1 rules make the FAIR files' undeclared linear units formally meters; incoming inspection explicitly uses millimeters. A plausible millimeter repair is therefore a separately labeled hypothesis, never an automatic correction. On fresh Box data, **76 values** include an observed violation of the receiving interval even under that hypothesis. On Cover, **400 selected values** satisfy the numerical interval, but the source records do not establish coverage of the required physical family. A permutation-invariant record result does not locate a hole for rework or release a part.

Evidence, including failures

Experiment	Finding	What it does not establish
Original frozen instrument	**5/8** primary scoped audits accepted	Population reliability or physical acceptance
Original fresh recipient	**1/1** evidence-acquisition audit accepted after actual byte transfer	Live ENS identity or production manufacturing authority
Prospective repair, separate version 0.1.1	**4/4** new requests accepted	Replacement of the original 5/8; independent physical replication
Cheap same-model trial	**U 4/4; E 4/4; E+ 3/4** under strict criteria	A measured expert or retained-method premium
Ordinary library with identical content/permissions	Same semantic result	Graph-specific value

The first instrument failed three actionability requirements because two repair requests had no affected-obligation bindings. A new version corrected this on Plate after new requests had already been frozen. Original failures and bytes remain intact. The new tests use fresh missions on the same historical data.

E+'s one hosted failure is a component-ID/serial inconsistency. A lenient reading would give it 4/4, still no gain over E. Calls were 7/6/8, with equal eight-call and 1,600-word ceilings. Complete tokens, billing, active latency and customer effort remain unknown. Lower call count alone is not a complete cost saving.

The transfer moved a **43,242-byte** bundle with the complete custom application; large public data and installed native dependencies were separately identified. All **35** study transfer controls passed, along with **29** inherited lifecycle controls. Expired/revoked grants, changed sources and altered missions were refused. The production capsule correctly retained its unsupported-domain boundary. Six cached integrity cases becoming zero repeated cases is not a reduction in total customer work.

Economics: positive modeled delivery is not incremental Alpha

No customer sale or matched engineering-service quote was obtained. Human-service prices are explicitly estimated from team hours and rates. Market needs, values, quality, costs and payment events are assumptions. **Actual collections and earned revenue are zero; complete actual research cost is unknown.**

The frozen model searches price, buyer subset and development terms against outside options, rights, quality, capacity and cash. It does not impose a fixed discount, investment count, coalition or training budget. Its central 80 assumed horizons yield:

Configuration	Feasible / 80	Mean, USD
Expert, standard	74	2,324.09
Expert, arranged	75	2,476.92
Retained, standard	27	492.72
Retained, arranged	38	771.24

Across eleven declared stresses, the combined package creates **zero new feasibility crossings against the expert standard baseline**. All central combined-feasible cases already qualify under that alternative. A prospective six-horizon retention campaign estimates USD2,220.94 benefit against USD2,950 additional construction, retention and transfer cost, and is rejected. Forced counterfactual construction ends at mean incremental value **−USD1,049.79**; it is not recovered profit.

The financing model assumes acceptance of its development/failure-record milestone. An independent later audit stresses rejection of that milestone and full advance refund after committed work. It identifies **114 of 1,889** original selected commitments with insufficient opening cash under that additional contingency. These are conditional exposures, not observed insolvencies. The original numbers remain unchanged, with their narrower assumption explicit.

An operating blueprint

Offer: a precisely scoped source-comparability and evidence-acquisition assessment. Price new measurements, authoritative corrections and release support separately. Determine the recipient's real configured outside option.

Deliver: automate inventory, literal checks, geometry traces and reproducible witnesses; independently verify material claims; ask for the smallest missing item; measure actual corrections, customer interaction and support.

Retain: generic executable readers, proof policies and adapters only when rights, scope, fresh acceptance and expected reuse justify additional costs. Keep customer secrets, restricted IP and protected answers outside the package.

Finance: state milestone outcomes, private adaptations, risk, refunds, responsibilities and collection dates. An advance is not earned surplus. Participation and liquidity are separate tests. Preserve cash when a proposal fails either one.

Share: use the existing protected network, eligible membership and current mission-specific authority. Access does not confer unrestricted execution or independent possession of a hosted specialist. No production scope was changed.

Next decisive evidence: unrelated industrial packets; actual human/integrator quotes or trials; buyer-specific rights and source corrections; complete costs; accepted contracts; and authorized production qualification. No Gym campaign or specialist ASI claim is warranted by the current evidence.

What this edition leaves behind

The enterprise gains an executable evidence instrument, a useful fresh recipient handoff, an honest search history, a repaired actionability defect and a tested stop decision. It gains a sharper account of what must be true before capability becomes opportunity. It does not gain invented customers, retrospective profit or automatic permission.

The full paper, code, source notices, frozen results, evidence companion and English/French demonstration accompany this brief. Research and writing used AI assistance; role-separated AI review is not independent human peer review.