AGI ALPHA FLYWHEEL — When Capability Creates Opportunity

Transformational Intelligence, Experience-Trained Specialists and the Expanding Frontier of Profitable Work

Vincent Boucher — President of MONTREAL.AI

EXECUTIVE BRIEF · 27 OCTOBER 2026

Decision

Retain the executable audit and evidence as research assets. Do not claim a newly feasible ecological business, an added capability advantage or commercial compounding. A corrected version and a real beneficiary test must precede general service admission.

What was learned

The study examined six serious opportunities before constructing one instrument. Native-seed purchases and substitutions can obscure original ecological requirements. NASEM documents that concern; established mix design, forward contracts and seed qualification provide strong alternatives. The research rejected those methods as flagship inventions. [1]

The instrument preserves the charter, distinguishes source coverage from qualified lot coverage, and reports unresolved identity, provenance, quantity, timing and seasonal/host obligations. Identical purchases with two different charters—flowering versus nectar evidence—produce different unmet requirements. This witnesses information loss, not autonomous intent discovery. Lots and commercial terms are fictional; the 21-taxon catalog contains attributed source facts. [2]

The development report correctly retains two blocked lot rows even when seasonal counts are sufficient. The actual frozen compiler then produces correct reports on **24 fresh cases**: eight positive documentary screens, thirteen valid negative reports and three invalid-input rejections. A distinct recipient agent uses its unchanged source on dry-context work and then wet-context work with a monarch-host requirement. [3]

Finite success; broader defects

Five selected post-trial diagnostics expose incomplete validation: omitted price, unsupported host, omitted source, empty lot identifier and a missing nectar field. The core remains frozen. The 24/24 result does not establish general contract compliance or production readiness. [4]

Competent baselines change the conclusion

Each workflow had the same contract, facts, tools and 480-second cap per packet. U received a reasonable direct request; E an explicit expert procedure; E+ that procedure plus the retained compiler. U and E could build and retain their own tools. No external scoring feedback was supplied before fresh work. [5]

Configuration	First packet	Fresh packet	Correct reports
U: direct use	87.00 s	29.00 s	24/24
E: expert procedure	107.80 s	44.36 s	24/24
E+: expert + artifact	82.66 s	41.69 s	24/24

E+ is faster than E in these executions; U is fastest overall. These first-command-to-verification intervals come from one actor per condition. No new task becomes feasible under the observed quality/deadline conditions. Bills and genuine recipient effort are unknown.

Economics before expansion

Actual sales and collections are zero. Complete actual research cost and profit are unknown. Unpaid research cannot be described as customer-funded learning. The executable accounting model keeps A, complete paid delivery; B, additional retention; and C, optional development disjoint. It distinguishes deposits, liabilities, earned fees, receivables, collections, refunds, reserves and available funds. [6]

H remains the primary commercial reference, with an estimated three hours at USD 100/hour rather than a quote. The equal-three-hour allocation is an *uncalibrated null control*, not automated delivery cost or required human review. Its USD 400 offer illustrates accounting. E+ spending gives expected contribution USD -73.14 and needs USD 457 against assumed USD 400 funds. Hypothetical joint construction gives USD -742 and needs USD 7,320 against USD 4,000. Actual evidence and participants are missing, so that project stops earlier. [6]

A transparent retention boundary illustrates the decision without inventing gain. At assumed fixed incremental spending of USD 100 and USD 2 extra serving cost per reuse, twenty-five qualifying reuses require **more than USD 6 gross benefit per reuse**. Neither that benefit nor the demand is measured. A profitable service could still be offered without investing in retention. [6]

Operating blueprint

1. **Offer a precise audit when a beneficiary values the decision.** Preserve the charter; compare with the recipient's competent process and existing tools. Use available solutions when they suffice.
2. **Repair and qualify before general live service.** Freeze a corrected version and test the five failure classes prospectively. Automate intake, checking and packaging; measure genuine residual effort.
3. **Fix quality, deadline, effort, price and payment conditions.** Seek a matched human quote or disclose estimated scope/rates. Let the buyer decline, defer or choose alternatives.
4. **Retain selectively; authorize separately.** Preserve version, dependencies, rights, provenance and scope. Recheck each mission. Exclude protected answers. Identity does not prove fitness or permission.
5. **Build technical evidence, participation and financing together.** Physical work needs source rights, viable material, actual participants, delivery evidence, failure allocation and a funded cash schedule.

Native interfaces pass 18 local source-inspection checks; they do not authenticate live `*.agent.agi.eth` agents or admit ecology. No Gym campaign or specialist ASI result is claimed. A separately frozen corrected investment selector passes six fresh hypothetical decision sets, choosing feasible alternatives before ranking contribution; this is policy validation, not actual profit. [7, 8]

Next decisive evidence

A real beneficiary must accept useful work against its own competent baseline, with complete costs and genuine willingness to pay. Fresh work must then justify the extra retention cost. Until those links hold, the deliverable is a reproducible instrument, its disclosed limitations and a defensible stopping decision.

Evidence keys. [1] `DISCOVERY.md`; NASEM, https://www.nationalacademies.org/read/26618/chapter/9. [2] `results/information_loss_witness.json`, `data/catalog.json`. [3] `METHOD_FREEZE.json`, `results/fresh_confirmation.json`. [4] `REVIEW.md`, `review/diagnostics.json`. [5] `results/workflow_assay.json`, `workflow_evidence/`. [6] `economics/results/report.json`, `economics/NULL_CONTROL_NOTE.md`, `results/retention_boundary.json`. [7] `results/native_integration.json`. [8] `POLICY_ADDENDUM.md`.

Provenance. AI-assisted research, construction and checking under the stated author direction; separate same-session roles, not external institutional validation. Earlier papers and evidence are preserved.